

Vivid Runtime Architecture

A technical reference for the Vivid audiovisual graph runtime.


1. Introduction

Vivid is a real-time audiovisual graph engine. Users build directed graphs of **operators** — nodes that generate or transform signals, audio, GPU textures, and data — and the runtime evaluates them continuously at two independent rates: a **frame cadence** (~60 Hz, main thread) for control logic and GPU rendering, and an **audio cadence** (48 kHz, dedicated thread) for sample-accurate sound processing.

The runtime's core responsibilities are:

- **Compile** a user-authored graph into an efficient, pre-allocated execution plan.
- **Execute** frame-rate and audio-rate operators in topological order on their respective threads.
- **Synchronize** data between cadences using lock-free double-buffered snapshots.
- **Hot-reload** operator code from recompiled dynamic libraries without stopping playback.

Vivid Runtime Architecture



Architecture Overview

2. Architecture Overview

The runtime is composed of several cooperating subsystems:

Component	Thread	Role
RuntimeCore	Main	Top-level orchestrator. Owns the compiled graph, cadence bridge, and frame executor. Drives the tick cycle.

Component	Thread	Role
GraphCompiler	Main (once)	Static factory that transforms a <code>Graph</code> into a <code>CompiledGraph</code> .
CompiledGraph	Both (read-only)	The execution plan: pre-allocated node state, classified edges, and topologically sorted execution orders.
FrameExecutor	Main	Iterates <code>frame_order</code> , propagates direct edges, calls <code>process_frame()</code> and <code>process_gpu()</code> .
AudioExecutor	Audio	Iterates <code>audio_order</code> inside the miniaudio callback, processes 256-sample buffers at 48 kHz.
CadenceBridge	Both	Double-buffered snapshot bridge. Transfers parameters frame-to-audio and analysis audio-to-frame without locks.
OperatorRegistry	Main	Discovers, lazy-loads, and hot-reloads operator dynamic libraries.
GpuContext	Main	Manages the WebGPU device, queue, and surface.
HotReloader	Background	Watches source files and runs background cmake builds.

3. The Graph Model

A Vivid graph is a JSON document containing nodes, connections, and metadata.

NodeDef

Each node declares an operator type and its parameter values:

```
{
  "osc": {
    "type": "Oscillator",
    "params": { "frequency": 440.0, "amplitude": 0.8, "waveform": 0 }
  }
}
```

Key fields: `type` (operator name), `params` (float values), `string_params` (file paths, text), `cadence_override` (Auto / Frame / Audio), `tex_width` / `tex_height` (per-node GPU resolution).

ConnectionDef

Connections wire an output port of one node to an input port of another:

```
{ "from": "osc/output", "to": "gain/input" }
```

Connections may carry an optional **value remap**: `from_min`, `from_max`, `to_min`, `to_max`, and a `clamp` flag. This linearly rescales the value as it crosses the edge.


Port Types

Type	ID	Description
SIGNAL	0	Scalar float (frame) or 1-channel buffer (audio)
AUDIO	1	Multi-channel sample buffer
SPREAD	2	Variable-length float array
STRING	3	UTF-8 text
STRING_SPREAD	4	Variable-length string array
TEXTURE	5	WGPUTextureView (GPU only)

Custom port types can be registered at runtime via the port type registry, identified by a `stable_type_id`.

4. Graph Compilation

The `GraphCompiler` transforms a `Graph` plus an `OperatorRegistry` into a `CompiledGraph` — a fully pre-allocated, ready-to-execute representation.



Graph Compilation Pipeline

Compilation Steps

1. **Resolve descriptors.** For each node, look up the `VividOperatorDescriptor` from the registry. This provides port definitions, parameter metadata, execution environment, and cadence capability.
2. **Assign cadences.** Each node is assigned `Cadence::Frame` or `Cadence::Audio`:
 - If the node has `cadence_override == Audio` and the operator is `AUDIO_CAPABLE` → Audio.
 - If `cadence_override == Frame` → Frame.
 - If the descriptor declares `has_process_audio` → Audio.
 - Otherwise → Frame.
3. **Classify edges.** Each connection becomes a `CompiledEdge` with a transport type:
 - **Direct** — both endpoints share the same cadence. The executor copies data during its pass.
 - **Snapshot** — endpoints are on different cadences. Data flows through the `CadenceBridge`.

4. **Topological sort.** Kahn's algorithm produces two independent execution orders:

- `frame_order` — all Frame and GPU nodes.
- `audio_order` — all Audio nodes.

5. **Pre-allocate state.** All buffers, parameter arrays, spread storage, string arrays, GPU textures, and audio buffers are allocated up front. No heap allocation occurs during execution.


CompiledGraph Structure

```
CompiledGraph
├─ nodes[]           ─ CompiledNode instances (all state)
├─ edges[]           ─ CompiledEdge instances
├─ frame_order[]      ─ topological indices for frame nodes
├─ audio_order[]      ─ topological indices for audio nodes
├─ frame_direct_edges[] ─ edge indices within frame cadence
├─ audio_direct_edges[] ─ edge indices within audio cadence
├─ frame_to_audio_edges[] ─ snapshot edge indices
└─ audio_to_frame_edges[] ─ snapshot edge indices
```

5. Dual-Cadence Execution

The two cadences run on independent threads at vastly different rates. They never block each other.

Dual-Cadence Tick Sequence



Dual-Cadence Tick Sequence

Frame Cadence (~60 Hz, Main Thread)

Each frame follows a three-phase pattern:

1. **pre_tick_audio_sync()** — Pull the latest `AnalysisSnapshot` from the audio thread. Inject RMS, peak, waveform, and scalar outputs into frame-rate nodes.
2. **FrameExecutor::tick()** — Iterate `frame_order`. For each node:
 - Propagate values along `frame_direct_edges` (copy, remap, spread merge).
 - Call `process_frame()` (or `process_gpu()` for GPU nodes).
 - Apply skip logic: only process if time-dependent, dirty, or upstream changed.
3. **post_tick_audio_sync()** — Snapshot frame outputs into the `CadenceBridge` for the audio thread to consume.

Audio Cadence (~48 kHz = ~188 callbacks/sec, Audio Thread)

The miniaudio callback fires approximately 188 times per second (48000 / 256):

1. Read the active `ParamSnapshot` (lock-free atomic acquire).
2. Iterate `audio_order`. For each node:
 - Propagate audio buffers along `audio_direct_edges`.
 - Handle auto-duplication for mono operators in multi-channel chains.
 - Call `process_audio()` with a `VividAudioContext`.
3. Compute per-node analysis (RMS, peak, waveform ring buffer).
4. Publish the `AnalysisSnapshot` (lock-free atomic release).


Timing Relationship

At 60 fps and 48 kHz with 256-sample buffers, approximately **3 audio callbacks** execute per visual frame. The two threads are decoupled — if the frame rate drops, audio continues uninterrupted.

6. The Cadence Bridge

The `CadenceBridge` is the sole communication channel between the frame and audio threads. It uses **double-buffered snapshots** with atomic index swaps — no mutexes touch the audio thread.

CadenceBridge: Double-Buffered Snapshots



Cadence Bridge

Frame → Audio: ParamSnapshot

Contains everything the audio thread needs from the frame world:

Field	Content
<code>node_params</code>	Parameter values per audio node
<code>float_input_values</code>	CV/signal inputs from frame-rate outputs
<code>spread_inputs</code>	Spread data crossing cadence boundary
<code>input_string_values</code>	String inputs (file paths, text)
<code>custom_inputs</code>	Custom-type port data
<code>solo_active_set</code>	Which nodes are active under solo mode

Write path (main thread): `push_to_audio()` writes to the *inactive* buffer, then atomically swaps `param_active_` with release semantics.

Read path (audio thread): `active_params()` reads `param_active_` with acquire semantics, returning the latest complete snapshot.

Audio → Frame: AnalysisSnapshot

Contains audio-thread outputs for visualization and frame-rate modulation:

Field	Content
<code>rms</code> , <code>peak</code>	Per-node level meters
<code>waveform</code>	1024-sample ring buffer per node
<code>float_outputs</code>	Scalar signal outputs for frame-rate consumption
<code>spread_outputs</code>	Spread data from audio nodes
<code>errored</code> , <code>error_msgs</code>	Error state (fixed-size, no heap)

The same double-buffer / atomic-swap pattern applies in the reverse direction.

7. Frame Executor

The `FrameExecutor` processes all frame-rate and GPU nodes on the main thread.

Tick Algorithm

```

for node in frame_order:
    if solo_active and node not in solo_set: skip

    zero input values

    for edge in frame_direct_edges targeting this node:
        copy output → input (with remap if configured)
        propagate spreads, strings, file params, custom ports

    build VividFrameContext with time, params, inputs, outputs, spreads
    call process_frame(instance, ctx)

    if GPU node:
        call process_gpu(instance, gpu_ctx)
        invoke PostNodeFn callback

```

Skip Logic

To avoid redundant computation on static graphs, the executor skips a node unless:

- It is **time-dependent** (reads `ctx->time`).

- It is a **root node** (no upstream connections).
- Any **upstream node** was processed this tick.
- The node is **dirty** (marked by bridge sync, API call, or hot-reload).

GPU Management

GPU nodes receive a `VividGpuContext` with the WebGPU device, queue, and command encoder. The executor manages per-node texture allocation, finds the GPU sink (`video_out`), and retrieves the final output texture for display.

8. Audio Executor

The `AudioExecutor` runs on the dedicated audio thread, processing 256-sample buffers at 48 kHz via miniaudio.

Audio Callback Flow

```
audio_callback(output_buffer, frame_count):
    snap = bridge.active_params()           // lock-free read

    apply snapshot params to audio nodes

    for node in audio_order:
        if solo_active and not in solo_set: mute

        propagate audio_direct_edges (buffer routing, channel negotiation)

        if auto_dup_group:
            deinterleave → per-channel mono buffers
            process each channel instance independently
            interleave → multi-channel output
        else:
            process_audio(instance, ctx)

    extract scalar float outputs for frame bridge

    compute RMS, peak, waveform per node
    bridge.publish_analysis()               // lock-free write
```

Channel Negotiation

Each audio port declares a channel count (0 = auto, 1 = mono, 2 = stereo, etc.). Edges negotiate channel counts between source and destination. When a mono operator appears in a stereo chain, the runtime uses **auto-duplication**.

Auto-Duplication

When a mono operator (1-in, 1-out) receives multi-channel input:

1. The runtime clones the operator instance for each channel.
2. Multi-channel buffers are deinterleaved into per-channel mono buffers.
3. Each instance processes its channel independently.
4. Outputs are interleaved back to multi-channel format.

This is transparent to the operator — it always sees mono buffers.


Diagnostics

The audio executor tracks: - **Underrun count** — atomic counter incremented on buffer underruns. - **Audio load** — ratio of processing time to buffer duration. - **RecordingTap** — lock-free ring buffer (960,000 samples = 10 sec stereo) for mix recording.

9. Operator System

Operators are the computational units of a Vivid graph. They are implemented as C++ structs compiled to dynamic libraries.

Operator Lifecycle



Operator Lifecycle

Defining an Operator

```
struct MyFilter : vivid::OperatorBase, vivid::FrameProcessable {
    static constexpr const char* kName = "MyFilter";
    static constexpr bool kTimeDependent = false;

    vivid::Param<float> cutoff{"cutoff", 1000.0f, 20.0f, 20000.0f};
    vivid::Param<int> mode{"mode", 0, 0, 3};

    void collect_params(std::vector<vivid::ParamBase*>& out) override {
        out.push_back(&cutoff);
        out.push_back(&mode);
    }

    void collect_ports(std::vector<VividPortDescriptor>& out) override {
        out.push_back({"in", VIVID_PORT_SIGNAL, VIVID_PORT_INPUT});
        out.push_back({"out", VIVID_PORT_SIGNAL, VIVID_PORT_OUTPUT});
    }

    void process_frame(const VividFrameContext* ctx) override {
        float c = ctx->param_values[0];
        ctx->output_values[0] = ctx->input_values[0] * c;
    }
};

VIVID_REGISTER(MyFilter)
```

Capability Interfaces

An operator opts into execution environments by inheriting capability interfaces:

Interface	Method	Environment
FrameProcessable	process_frame(VividFrameContext*)	Main thread, ~60 Hz
AudioProcessable	process_audio(VividAudioContext*)	Audio thread, 48 kHz
GpuProcessable	process_gpu(VividGpuContext*)	Main thread, WebGPU

An operator may implement multiple interfaces. The `cadence_capability` flag (`FRAME_ONLY` or `AUDIO_CAPABLE`) determines whether the operator can be promoted to audio-rate execution.

The VIVID_REGISTER Macro

This macro generates the `extern "C"` entry points that the runtime loads via `dlopen` :

- `vivid_descriptor()` — returns `VividOperatorDescriptor` with all metadata.
- `vivid_create()` / `vivid_destroy()` — instance lifecycle.

- `vivid_process_frame()`, `vivid_process_audio()`, `vivid_process_gpu()` — execution dispatch.
- `vivid_main_thread_update()` — optional per-frame hook on main thread.

VividOperatorDescriptor

The descriptor is the operator's complete metadata:

<code>name</code>	— type name (e.g., "Oscillator")
<code>param_count, params[]</code>	— parameter descriptors (name, type, range, defaults)
<code>port_count, ports[]</code>	— port descriptors (name, type, direction, transport)
<code>time_dependent</code>	— whether the operator reads <code>ctx->time</code>
<code>execution_env</code>	— <code>VIVID_ENV_FRAME</code> / <code>AUDIO</code> / <code>GPU</code>
<code>cadence_capability</code>	— <code>FRAME_ONLY</code> or <code>AUDIO_CAPABLE</code>
<code>embedded_op_slots[]</code>	— modulation sub-operator slots


OperatorRegistry

The registry manages operator discovery and loading:

- `scan_deferred(dir)` — Probes each `.dylib` for metadata without full load.
- `load_for_graph(graph)` — Lazy-loads only the operators referenced by the graph.
- `find(type_name)` — Returns the `OperatorLoader` (may trigger lazy load).
- `register_builtin()` — Registers built-in operators (e.g., `audio_out`, `video_out`).
- `register_alias()` — Maps alternative names to canonical types.

10. Data Flow & Edge Types

Data Flow & Edge Types



Data Flow and Edge Types

Direct Edges

Same-cadence edges are **Direct**: the executor copies the output value from the source node into the input slot of the destination node during its iteration pass. This is a simple memory copy with no synchronization overhead.

For **SIGNAL** ports, a single float is copied. For **AUDIO** ports, the entire sample buffer pointer is routed. For **SPREAD** and **STRING** ports, the data is copied into pre-allocated staging buffers.

Snapshot Edges

Cross-cadence edges are **Snapshot**: data is written into the `CadenceBridge`'s double-buffered snapshot and read by the other thread on its next cycle. There is an inherent one-cycle latency (one frame period for frame-to-audio, one audio callback period for audio-to-frame).

Value Remap

Any edge can carry a remap transform:

```
normalized = (value - from_min) / (from_max - from_min)
result      = to_min + normalized * (to_max - to_min)
```

With an optional `clamp` to constrain the result to `[to_min, to_max]`. This allows connecting outputs with different value ranges (e.g., a 0–1 LFO driving a 20–20000 Hz frequency parameter).

Port Type Compatibility

- **SIGNAL** and **AUDIO** are cross-compatible: at audio cadence, a SIGNAL port becomes a 1-channel buffer.
- Control types (SIGNAL, SPREAD, STRING) are inter-compatible for flexible routing.
- **TEXTURE** ports connect only to other TEXTURE ports.
- **Custom** ports match by `stable_type_id`.

11. Hot Reload

Vivid supports hot-reloading operator code while the graph is running.

Pipeline

1. **FileWatcher** detects a source file change.
2. **HotReloader** queues a background `cmake --build` for the affected target.
3. The background compile thread produces a new `.dylib`.
4. Main thread polls `poll_ready()` and initiates the reload.

Two-Phase Reload

Reloading is a delicate operation because the old dylib must remain loaded while old instances are destroyed:

1. **Phase 1: Destroy old instances.** For all nodes of the reloading type, save current parameter values, then call `destroy_instance()` while the old loader is still valid.
2. **Phase 2: Swap and recreate.** Replace the dylib in the registry, create new instances via the new loader, and reconcile saved parameters (matching by name to handle added/removed params).

For audio operators, the `AudioEngine` coordinates a pause/resume around the reload to avoid accessing stale instances on the audio thread.

12. Appendix: Key Types Reference

VividFrameContext

Passed to frame-rate operators on the main thread.

Field	Type	Description
<code>time</code>	<code>double</code>	Absolute time in seconds
<code>delta_time</code>	<code>double</code>	Frame duration (~0.0167s at 60 Hz)
<code>frame</code>	<code>uint64_t</code>	Frame counter
<code>param_values</code>	<code>float*</code>	Parameter array, indexed by param order
<code>input_values</code>	<code>float*</code>	Input port values (SIGNAL)
<code>output_values</code>	<code>float*</code>	Output port values (SIGNAL) — write here
<code>input_spreads</code>	<code>VividSpreadPort*</code>	Spread input ports
<code>output_spreads</code>	<code>VividSpreadPort*</code>	Spread output ports
<code>input_string_values</code>	<code>const char**</code>	String input values
<code>output_string_values</code>	<code>const char**</code>	String output values
<code>file_param_values</code>	<code>const char**</code>	File/text parameter values
<code>input_state</code>	<code>const VividInputState*</code>	Keyboard/mouse input

VividAudioContext

Passed to audio-rate operators on the audio thread.

Field	Type	Description
<code>time</code>	<code>double</code>	Absolute time in seconds
<code>delta_time</code>	<code>double</code>	Chunk duration (buffer_size / sample_rate)
<code>frame</code>	<code>uint64_t</code>	Sample frame counter

Field	Type	Description
<code>param_values</code>	<code>float*</code>	Parameters (from ParamSnapshot)
<code>input_buffers</code>	<code>float**</code>	Planar audio: <code>[port][sample * channels]</code>
<code>output_buffers</code>	<code>float**</code>	Planar audio output — write here
<code>buffer_size</code>	<code>uint32_t</code>	Samples per buffer (256)
<code>sample_rate</code>	<code>uint32_t</code>	Sample rate (48000)
<code>input_channel_counts</code>	<code>const uint8_t*</code>	Per-port input channel count
<code>output_channel_counts</code>	<code>const uint8_t*</code>	Per-port output channel count
<code>input_float_values</code>	<code>float*</code>	CV inputs from frame cadence
<code>output_float_values</code>	<code>float*</code>	Scalar outputs for frame cadence
<code>channel_index</code>	<code>uint8_t</code>	Channel index (for auto-dup groups)

CompiledNode (key fields)

Field	Description
<code>node_id</code>	Unique identifier from graph
<code>type_name</code>	Operator type name
<code>active_cadence</code>	<code>Cadence::Frame</code> or <code>Cadence::Audio</code>
<code>loader / instance</code>	Operator loader and live instance
<code>param_values[]</code>	Current parameter floats
<code>input_values[] / output_values[]</code>	SIGNAL port state
<code>input_spreads[] / output_spreads[]</code>	Spread port state
<code>audio</code>	<code>AudioNodeState*</code> — buffers, channels (audio nodes only)
<code>gpu</code>	<code>GpuNodeState*</code> — textures, views (GPU nodes only)

CompiledEdge (key fields)

Field	Description
<code>from_node / to_node</code>	Source and destination node indices
<code>from_port / to_port</code>	Port ordinals

Field	Description
<code>transport</code>	<code>Direct</code> or <code>Snapshot</code>
<code>data_type</code>	Port type (SIGNAL, AUDIO, SPREAD, etc.)
<code>from_channels</code> / <code>to_channels</code>	Audio channel negotiation
<code>from_min/max</code> , <code>to_min/max</code> , <code>clamp</code>	Value remap parameters